

* Please Read Observations Bottom Of The Notebook *

Link Of Dataset : <https://www.kaggle.com/puneet6060/intel-image-classification>
(<https://www.kaggle.com/puneet6060/intel-image-classification>)

Load Data From Kaggle and Extract From Zip

In [2]:

```
--header="Host: storage.googleapis.com" --header="User-Agent: Mozilla/5.0 (X11; Lin
```

```
--2020-05-02 09:22:45-- https://storage.googleapis.com/kaggle-data-sets/111880/269359/bundle/archive.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1588426820&Signature=S7eh8bw5uxQpFniUMdVUiJCJuigSUHszRQexS0YoWAEP0Y1VSKYcZWAdZCazLBT5Eg6zX0HZwK06GPvoWGPZUdn9ZAXSGN%2FBIq9J3yknLwCw%2FMc9%2Blm5IAiGGimbnvVaNtVGATEYyzj6yRRI6tImDK80cXgbx03FoRQGtNL1gGKlGplCchRb44IFaTQFp2o7By0mpCeVXgFxtkuef7C3Fv3JbEAdNjo8vjh7oniNVXYlBY8maSy0ZU28QTnHoVZ%2B5fs7LK6KXwd3jLM3QFABvf07tyDY0fjjXmc56%2Bood%2F5J%2FHZGLelBtsepBf48Zj%2FmbIH3ST1x9Mpj7yq8pkxS5A%3D%3D&response-content-disposition=attachment%3B+filename%3Dintel-image-classification.zip (https://storage.googleapis.com/kaggle-data-sets/111880/269359/bundle/archive.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1588426820&Signature=S7eh8bw5uxQpFniUMdVUiJCJuigSUHszRQexS0YoWAEP0Y1VSKYcZWAdZCazLBT5Eg6zX0HZwK06GPvoWGPZUdn9ZAXSGN%2FBIq9J3yknLwCw%2FMc9%2Blm5IAiGGimbnvVaNtVGATEYyzj6yRRI6tImDK80cXgbx03FoRQGtNL1gGKlGplCchRb44IFaTQFp2o7By0mpCeVXgFxtkuef7C3Fv3JbEAdNjo8vjh7oniNVXYlBY8maSy0ZU28QTnHoVZ%2B5fs7LK6KXwd3jLM3QFABvf07tyDY0fjjXmc56%2Bood%2F5J%2FHZGLelBtsepBf48Zj%2FmbIH3ST1x9Mpj7yq8pkxS5A%3D%3D&response-content-disposition=attachment%3B+filename%3Dintel-image-classification.zip)
```

```
Resolving storage.googleapis.com (storage.googleapis.com)... 108.177.98.128, 2607:f8b0:400e:c00::80
```

```
Connecting to storage.googleapis.com (storage.googleapis.com)|108.177.98.128|:443... connected.
```

```
HTTP request sent, awaiting response... 200 OK
```

```
Length: 363152213 (346M) [application/zip]
```

```
Saving to: 'intel-image-classification.zip'
```

```
intel-image-classif 100%[=====>] 346.33M 77.3MB/s in 4.5s
```

```
2020-05-02 09:22:49 (77.3 MB/s) - 'intel-image-classification.zip' saved [363152213/363152213]
```

In [3]:

```
!ls
```

```
intel-image-classification.zip sample_data
```

In [4]:

```
!unzip intel-image-classification.zip
```

Streaming output truncated to the last 5000 lines.

```
inflating: seg_train/seg_train/mountain/7537.jpg
inflating: seg_train/seg_train/mountain/7539.jpg
inflating: seg_train/seg_train/mountain/7551.jpg
inflating: seg_train/seg_train/mountain/7560.jpg
inflating: seg_train/seg_train/mountain/7565.jpg
inflating: seg_train/seg_train/mountain/7578.jpg
inflating: seg_train/seg_train/mountain/7581.jpg
inflating: seg_train/seg_train/mountain/7586.jpg
inflating: seg_train/seg_train/mountain/7647.jpg
inflating: seg_train/seg_train/mountain/7652.jpg
inflating: seg_train/seg_train/mountain/7654.jpg
inflating: seg_train/seg_train/mountain/7662.jpg
inflating: seg_train/seg_train/mountain/767.jpg
inflating: seg_train/seg_train/mountain/7672.jpg
inflating: seg_train/seg_train/mountain/7679.jpg
inflating: seg_train/seg_train/mountain/7681.jpg
inflating: seg_train/seg_train/mountain/7693.jpg
inflating: seg_train/seg_train/mountain/7695.jpg
inflating: seg_train/seg_train/mountain/7699.jpg
```

Import Libraries

In [1]:

```
import os
import shutil
from tqdm import tqdm

import cv2
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.utils import shuffle
from sklearn import metrics
from sklearn.metrics import confusion_matrix

%tensorflow_version 2.x
import tensorflow as tf
from tensorflow.keras import applications
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dropout, Flatten, Dense, Conv2D, MaxPooling2D,
from tensorflow.keras import optimizers
from tensorflow.keras.layers import Input
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras import optimizers

# tf.compat.v1.logging.set_verbosity(tf.compat.v1.logging.ERROR)
```

```
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:1
9: FutureWarning: pandas.util.testing is deprecated. Use the functions
in the public API at pandas.testing instead.
import pandas.util.testing as tm
```

In [4]:

```
sns.set_style("whitegrid") # For Graph Background White on Colab
```

Define Image Data Generator

In [14]:

```
train_datagen = ImageDataGenerator(rescale = 1./255, rotation_range=30, width_shift
                                brightness_range=None, shear_range=0.0, zoom_ran
                                channel_shift_range=0.0, fill_mode='nearest', ho
                                vertical_flip=False, validation_split=0.0, dtype=
```

In [15]:

```
test_datagen = ImageDataGenerator(rescale = 1./255, rotation_range=30, width_shift
                                brightness_range=None, shear_range=0.0, zoom_ran
                                channel_shift_range=0.0, fill_mode='nearest', hor
                                vertical_flip=False, validation_split=0.0, dtype=
```

Generate Train and Test Image Data

In [20]:

```
train_directory = './seg_train/seg_train/'
test_directory = './seg_test/seg_test/'

n_batch_size = 256
n_epochs = 50
```

In [6]:

```
train_generator = train_datagen.flow_from_directory(train_directory,
                                                    target_size=(150, 150),
                                                    batch_size = n_batch_size,
                                                    class_mode = 'categorical')
```

Found 14034 images belonging to 6 classes.

In [53]:

```
test_generator = test_datagen.flow_from_directory(test_directory,
                                                  target_size=(150, 150),
                                                  batch_size = n_batch_size,
                                                  class_mode = 'categorical')
```

Found 3000 images belonging to 6 classes.


In [25]:

```
train_labels = train_generator.classes
test_labels = test_generator.classes
class_names = list(test_generator.class_indices.keys())
```

Graphical Representation Of Data

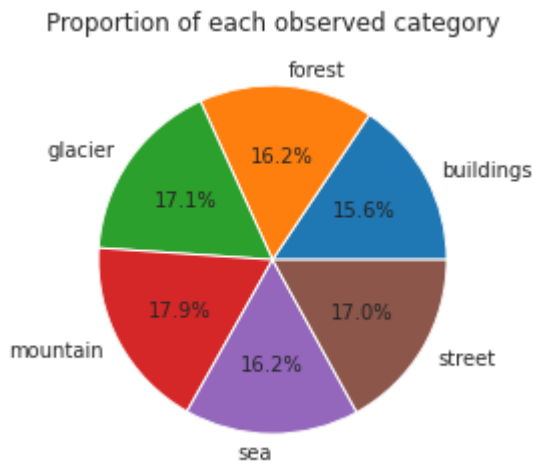
In [9]:

```
_, train_counts = np.unique(train_labels, return_counts=True)
_, test_counts = np.unique(test_labels, return_counts=True)
pd.DataFrame({'train': train_counts, 'test': test_counts}, index=class_names).plot.
plt.show()
```



In [10]:

```
plt.pie(train_counts, explode=(0, 0, 0, 0, 0, 0) , labels=class_names, autopct='%1.
#plt.axis('equal')
plt.title('Proportion of each observed category')
plt.show()
```



Display Some Train Images

In [31]:

```
def display_train_data(class_names, images, labels, n_rows=5 , n_cols=5 , fig_width
    """
    Display n images from the images array with its corresponding labels
    """
    n_images = n_rows * n_cols


    fig = plt.figure(figsize=(fig_width, fig_height))
    fig.suptitle("First " + str(n_images) + " Images of the Dataset", fontsize=16)
    for i in range(n_images):
        plt.subplot(n_rows,n_cols,i+1)
        plt.xticks([])
        plt.yticks([])
        plt.grid(False)
        plt.imshow(images[i], cmap=plt.cm.binary)
        plt.xlabel(class_names[np.where(labels[i] == 1)[0][0]])
    plt.show()
```

In [13]:

```
images, labels = train_generator.next()
```

```
display_train_data(class_names, images, labels, n_rows=5 , n_cols=5 , fig_width=10,
```


First 25 Images of the Dataset



sea



buildings



mountain



glacier



sea




mountain



mountain




glacier



glacier



sea



forest




glacier




glacier




mountain




forest



mountain



sea




glacier




street




buildings




glacier




mountain



glacier



glacier



sea

Load VGG16 as Base Model

In [18]:

```
vgg_model = applications.VGG16(weights='imagenet', include_top=False, input_shape=(  
print('Model Loaded.')
```

Model Loaded.

Create Custom Layers Top of on VGG16 Model

In [40]:

```
vgg_model_output = vgg_model.output

top_model = Conv2D(filters=32, kernel_size=(3, 3), activation='relu', padding="same")
top_model = Conv2D(filters=64, kernel_size=(3, 3), activation='relu', padding="same")
top_model = MaxPooling2D(pool_size=(2, 2), strides=(2, 2))(top_model)
top_model = Dropout(0.25)(top_model)

top_model = Flatten()(top_model)

top_model = Dense(256, activation='relu')(top_model)
top_model = Dropout(0.5)(top_model)

top_model = Dense(512, activation='relu')(top_model)
top_model = Dropout(0.5)(top_model)

top_model = Dense(6, activation='softmax')(top_model)

custom_model = Model(inputs=vgg_model.input, outputs=top_model)
```

Active Last 14 Layers

In [20]:

```

for layer in custom_model.layers[:-14]:
    layer.trainable = False

for layer in custom_model.layers:
    print(layer, layer.trainable)

```

```

<tensorflow.python.keras.engine.input_layer.InputLayer object at 0x7f5ae241f208> False
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae24390f0> False
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae2439438> False
<tensorflow.python.keras.layers.pooling.MaxPooling2D object at 0x7f5ae2485b00> False
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae2485e10> False
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae246bba8> False
<tensorflow.python.keras.layers.pooling.MaxPooling2D object at 0x7f5ae2e613c8> False
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae232ac50> False
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae232add8> False
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae235d588> False
<tensorflow.python.keras.layers.pooling.MaxPooling2D object at 0x7f5ae23374a8> False
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae2356b70> False
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae23482e8> False
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae2341518> False
<tensorflow.python.keras.layers.pooling.MaxPooling2D object at 0x7f5ae2322390> False
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae2322240> True
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae233c588> True
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae232f748> True
<tensorflow.python.keras.layers.pooling.MaxPooling2D object at 0x7f5ae237a668> True
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae25df9e8> True
<tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f5ae25dfcc0> True
<tensorflow.python.keras.layers.pooling.MaxPooling2D object at 0x7f5ae25e3ba8> True
<tensorflow.python.keras.layers.core.Dropout object at 0x7f5ae25e3da0> True
<tensorflow.python.keras.layers.core.Flatten object at 0x7f5ae25edc18> True
<tensorflow.python.keras.layers.core.Dense object at 0x7f5ae25f3048> True
<tensorflow.python.keras.layers.core.Dropout object at 0x7f5ae25edfd0> True
<tensorflow.python.keras.layers.core.Dense object at 0x7f5ae24e86a0> True

```

```
rue
<tensorflow.python.keras.layers.core.Dropout object at 0x7f5ae24e8198>
True
<tensorflow.python.keras.layers.core.Dense object at 0x7f5ae24f43c8> T
rue
```

Compile Our Model

In [44]:

```
custom_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['a
```

Generate Directory for Save Checkpoints after every Epoch

In [0]:

```
# shutil.rmtree(checkpoint_dir)
```

In [47]:

```
# Save The Checkpoints

# Include the epoch in the file name (uses `str.format`)
cp_dir = 'saved_checkpoints/'
if not os.path.exists(cp_dir):
    os.mkdir(cp_dir)

checkpoint_path = cp_dir+"cp-{epoch:04d}.ckpt"
checkpoint_dir = os.path.dirname(checkpoint_path)
```

In [48]:

```
# Create a callback that saves the model's weights every 5 epochs
cp_callback = ModelCheckpoint(
    filepath=checkpoint_path,
    verbose=0,
    save_weights_only=True,
    save_freq=5)

# Save the weights using the `checkpoint_path` format
custom_model.save_weights(checkpoint_path.format(epoch=0))
```

Start Training

In [149]:

```
trained_history = custom_model.fit( train_generator, steps_per_epoch = len(train_ge
                                   epochs = n_epochs, use_multiprocessing=False, validation_
                                   workers = 1, validation_steps=len(test_generator), callba
```

```
Epoch 1/50
55/55 [=====] - 97s 2s/step - loss: 1.2514
- accuracy: 0.4326 - val_loss: 0.7622 - val_accuracy: 0.7023
Epoch 2/50
55/55 [=====] - 96s 2s/step - loss: 0.6823
- accuracy: 0.7586 - val_loss: 0.5522 - val_accuracy: 0.8063
Epoch 3/50
55/55 [=====] - 96s 2s/step - loss: 0.5234
- accuracy: 0.8269 - val_loss: 0.4429 - val_accuracy: 0.8390
Epoch 4/50
55/55 [=====] - 96s 2s/step - loss: 0.4408
- accuracy: 0.8516 - val_loss: 0.4131 - val_accuracy: 0.8590
Epoch 5/50
55/55 [=====] - 96s 2s/step - loss: 0.4461
- accuracy: 0.8508 - val_loss: 0.4140 - val_accuracy: 0.8570
Epoch 6/50
55/55 [=====] - 96s 2s/step - loss: 0.3999
- accuracy: 0.8684 - val_loss: 0.4361 - val_accuracy: 0.8567
Epoch 7/50
55/55 [=====] - 96s 2s/step - loss: 0.3700
- accuracy: 0.8700 - val_loss: 0.4361 - val_accuracy: 0.8567
```

In [0]:


Plot Graph to Compar Train & Test Accuracy and Loss

In [51]:

```
def plot_graph(train_history, train_acc, test_acc, title, x_label, y_label):
    plt.plot(train_history.history[train_acc], label='Train')
    plt.plot(train_history.history[test_acc], label='Test')
    plt.title(title)
    plt.ylabel(y_label)
    plt.xlabel(x_label)
    plt.legend(loc='lower left')
    plt.grid()
    plt.show()
```

In [150]:

```
title = 'Model Accuracy'  
x_label = 'Epoch'  
y_label = 'Accuracy'  
plot_graph(trained_history, 'accuracy', 'val_accuracy', title, x_label, y_label)
```




In [151]:

```

title = 'Model Loss'
x_label = 'Epoch'
y_label = 'Loss'
plot_graph(trained_history, 'loss', 'val_loss', title, x_label, y_label)

```



In [0]:

Load Weights which has Maximum Accuracy and Minimum Loss

In [24]:

```

# Get Latest Checkpoint from Saved Directory
# latest_cp = tf.train.latest_checkpoint(checkpoint_dir)
# print(latest_cp)

cp = checkpoint_dir+'/cp-0019.ckpt'

# Create a new model instance
# custom_model = create_model()

# Load the previously saved weights
custom_model.load_weights(cp)

# Re-evaluate the model
loss, acc = custom_model.evaluate(test_generator, steps=len(test_generator), verbose=1)
print("Restored Model, Accuracy: {:.2f}%".format(100*acc))

```

```

12/12 [=====] - 15s 1s/step - loss: 0.3190 -
accuracy: 0.8907
Restored Model, Accuracy: 89.07%

```

In [0]:

Generate Test Data For Evaluation

In [201]:

```
test_generator = test_datagen.flow_from_directory(test_directory,
                                                  target_size=(150, 150),
                                                  batch_size = n_batch_size,
                                                  class_mode = 'categorical',
                                                  shuffle = False)

# set 'shuffle = False' in test_data_generator when you use 'flow_from_directory'
```

Found 3000 images belonging to 6 classes.

Prediction For Test Data

In [35]:

```
test_generator.reset()
predictions = custom_model.predict(test_generator, steps=len(test_generator), verbose=1)

12/12 [=====] - 15s 1s/step
```

In [0]:

Represent The Match Labels and Mismatch Labels in the form of Confusion Matrix

In [203]:


```

pred_labels = np.argmax(predictions, axis = 1)
test_labels = test_generator.classes
class_names = list(test_generator.class_indices.keys())

CM = confusion_matrix(test_labels, pred_labels)

ax = plt.axes()
sns.heatmap(CM, annot=True, xticklabels=class_names, yticklabels=class_names, cmap=
ax.set_title('Confusion matrix')
plt.show()

```



Detaile Report Of Accuracy

In [169]:

```

report = metrics.classification_report(test_labels, pred_labels, target_names=class
print(report)

```

	precision	recall	f1-score	support
buildings	0.90	0.89	0.89	437
forest	0.97	0.98	0.98	474
glacier	0.84	0.86	0.85	553
mountain	0.87	0.82	0.84	525
sea	0.89	0.93	0.91	510
street	0.89	0.91	0.90	501
accuracy			0.89	3000
macro avg	0.89	0.89	0.89	3000
weighted avg	0.89	0.89	0.89	3000

Trained Model Apply on New (or Unseen) Data

Load New Data

In [53]:

```
IMAGE_SIZE = (150, 150)
```

In [54]:

```
def load_new_data(datasets):
    images = []

    print("\nLoading {}".format(datasets))
    # Iterate through each image in our folder
    for file in tqdm(os.listdir(datasets)):

        # Get the path name of the image
        img_path = os.path.join(datasets, file)

        # Open and resize the img
        image = cv2.imread(img_path)
        image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
        image = cv2.resize(image, IMAGE_SIZE)

        # Append the image and its corresponding label to the output
        images.append(image)

    images = np.array(images, dtype = 'float32')

    return images
```

In [65]:

```
new_images = load_new_data('./seg_pred/seg_pred/')
3%|| | 202/7301 [00:00<00:03, 2018.44it/s]
```

Loading ./seg_pred/seg_pred/

100%|██████████| 7301/7301 [00:03<00:00, 2014.12it/s]

In [55]:

```
# Scale the Data
new_images = new_images / 255.0
```

Predict New Data

In [56]:

```
predictions_new_data = custom_model.predict(new_images)
pred_labels_new_data = np.argmax(predictions_new_data, axis = 1)
```

Display Some Predicted Images with Labels

In [60]:


```
def display_data(class_names, images, labels, n_rows=5 , n_cols=5 , fig_width=10, f
    """
    Display n images from the images array with its corresponding labels
    """
    n_images = n_rows * n_cols

    fig = plt.figure(figsize=(fig_width, fig_height))
    fig.suptitle("First " + str(n_images) + " Images of the Dataset", fontsize=16)
    for i in range(n_images):
        plt.subplot(n_rows,n_cols,i+1)
        plt.xticks([])
        plt.yticks([])
        plt.grid(False)
        plt.imshow(images[i], cmap=plt.cm.binary)
        plt.xlabel(class_names[labels[i]])
    plt.show()
```

In [82]:

```
display_data(class_names, new_images, pred_labels_new_data, n_rows=7, n_cols=7 , fi
```

First 49 Images of the Dataset



Save the Predicted Images in their Class Folders

In [65]:

```
def classify_new_images(class_names, images, labels):  
  
    dirs = 'new_data_output'  
    if not os.path.exists(dirs):  
        os.mkdir(dirs)  
  
    for i in range(len(images)):  
  
        class_nm = class_names[labels[i]]  
  
        class_path = os.path.join(dirs, class_nm)  
  
        if not os.path.exists(class_path):  
            os.mkdir(class_path)  
  
        filename = class_path+'/'+class_nm+str(f"{i:04}")+'.jpg'  
        cv2.imwrite(filename, images[i])
```

In [66]:

```
classify_new_images(class_names, new_images*255.0 , pred_labels_new_data)
```

In [0]:

```
# shutil.rmtree('new_data_output')
```

Link of Trained Model

Download Trained Model: https://drive.google.com/open?id=1-2nL_Tt8mWIKZRv-u8P7tLjLYVeT0Lis
(https://drive.google.com/open?id=1-2nL_Tt8mWIKZRv-u8P7tLjLYVeT0Lis)

In []:

```
# Load Model After Download  
from tensorflow.keras.models import load_model  
  
custom_model = load_model('custom_model_89.h5')
```

* Observations *

1. The data is very complicated with very small size images.
2. I tried various pre-processing steps but because of complicated data and small size images I got 89% accuracy.
3. Also I tried with unseen data given by the kaggle competition. Then I got good results.
4. The mismatch class images are also similar to match class images. (For example, "glacier & sea", "glacier & mountain", "street & buildings" images are more similar to each other.)

5. If data have high quality images then I think accuracy increase by 5% to 7% of current accuracy. (i.e, accuracy would be 94% to 96%).

In []: